

DSTI DEEP LEARNING PROJECT – COMPUTER VISION

SKIN CANCER DETECTION



PLAN ON A PAGE

| | Problem Analysis | | Data Annotation and Labelling | | Model Evaluation and Iteration | |
|--------------------------------------|--|---|---|--|--------------------------------|--|
| | Week 1 | Week 2 | Week 3 | Week 4 | Week 5 | Week 6 |
| | | Data Set Collection | | Model Selection & Training | | Final Iteration & Deployment |
| Data Science | Model testing, full team | | | | | |
| | Models Review | Week 2 | Week 3 | Week 4 | Week 5 | Week 6 |
| | <ul style="list-style-type: none"> Review of Hugging face and other sources Top 5 recommendation | <ul style="list-style-type: none"> Data set review vs top 5 models Model reprioritisation | <ul style="list-style-type: none"> Initial Feature Selection & Engineering Model reprioritisation (if required) | <ul style="list-style-type: none"> Model training & evaluation – base data set Evaluation report Iteration strategy, e.g. what to try next | | <ul style="list-style-type: none"> Ready Model for Kaggle Notebook (Production) deployment |
| Data Engineer | Model testing, full team | | | | | |
| | Techniques Review | Week 2 | Week 3 | Week 4 | Week 5 | Week 6 |
| | <ul style="list-style-type: none"> Review of image manipulation techniques Top 5 recommendation | <ul style="list-style-type: none"> Data set review Storage selection Data Processing Data Cleansing | <ul style="list-style-type: none"> <u>Additional Data Sets</u> Data set review Storage selection Data Processing Data Cleansing | <ul style="list-style-type: none"> Ad hoc data engineering support for model training Pick-up DS tasks as needed | | <ul style="list-style-type: none"> Ready data for Kaggle Notebook (Production) deployment |
| Data Analysis & Platform Engineering | Model testing, full team | | | | | |
| | Platform Analysis | Week 2 | Week 3 | Week 4 | Week 5 | Week 6 |
| | <ul style="list-style-type: none"> Review of Kaggle Notebooks Review & recommend dev environment, CoLab, AWS, Local, etc | <ul style="list-style-type: none"> Platform configuration Identification of additional data sets | <ul style="list-style-type: none"> <u>Data Analysis & Exploration</u> Statistical Exploration (label count, etc.) Data Viz (pair plots, histograms, etc) | <ul style="list-style-type: none"> Model analysis and visualisation support 1st Draft support Kaggle competition material (viz, slides, report, etc.) | | <ul style="list-style-type: none"> Ready support material for Kaggle competition submission |



PLATFORM ANALYSIS



DEVELOPMENT ENVIRONMENTS

| Feature | Google Colab | Local Environment | Kaggle Notebooks |
|-------------------------|--------------------------|------------------------|----------------------------------|
| Cost | Free (with paid options) | One-time hardware cost | Free |
| Ease of Use | Very easy | Moderate | Very easy |
| GPU Access | Yes | Depends on hardware | Yes (limited) |
| Scalability | Limited | Limited by hardware | Limited |
| Integration | Google Drive | Full control | Kaggle datasets & competitions |
| Runtime Limits | 12 hours max | No limits | 9 hours/week (free tier) |
| Pre-installed Libraries | Some | Manual installation | Extensive data science libraries |
| Data Access | Flexible | Local storage | Direct access to Kaggle datasets |
| Collaboration | Good | Limited | Excellent for competitions |
| Persistence | Session-based | Persistent | Version control included |

DATA STORAGE OPTIONS

| Feature | GitHub | AWS S3 | Google Cloud Storage | Kaggle Datasets |
|-----------------|-----------------|---------------------|----------------------|-----------------|
| Cost | Free for public | ~\$0.023/GB/month | ~\$0.020/GB/month | Free up to 20GB |
| Performance | Average | Excellent | Excellent | Good |
| File Size Limit | 100MB per file | Virtually unlimited | Virtually unlimited | 20GB total |
| Integration | Version control | AWS services | Google services | Kaggle platform |
| Ease of Use | Easy | Moderate | Moderate | Very easy |
| Scalability | Limited | Highly scalable | Highly scalable | Limited |

KAGGLE VS. COLAB NOTEBOOKS

| | Kaggle Notebooks | Google Colab Notebooks |
|------------------|---|---|
| Primary Use | Data science competitions | General-purpose |
| Environment | Pre-configured | Requires more setup |
| Data Access | Direct access to Kaggle datasets | Flexible, requires setup |
| Collaboration | Competition-focused | General collaboration |
| Runtime | Limited | Up to 12 hours |
| Integration | Kaggle ecosystem | Google ecosystem |
| GPU quota | 30 hours per week (in the free tier) | Unlimited in theory, but with fair usage limits |
| GPU types | Usually Nvidia P100 | Varies, can include T4, P100, V100 |
| GPU assignment | Automatic GPUs assignment based on availability | Manual selection by the user |
| GPU Availability | Consistent performance with reliable GPU | Inconsistent: GPUs might not always be available, performance can vary between sessions |

SETTING UP KAGGLE PROJECT

| | Description | |
|---------------------------|-------------------------------|--|
| Create or Connect Dataset | Upload and configure data | |
| Create Notebook | Set up coding environment | |
| Configure Notebook | Access data, import libraries | |
| Develop Model | Code, train, and test model | |
| Create Submission | Format results | |
| Submit Results | Upload to competition | |

TRANSITIONING FROM KAGGLE TO COLAB/AWS

| Aspect | Change Required | Complexity | Key Considerations |
|--------------------|----------------------------|------------|---------------------------------------|
| Code Transfer | Minor adjustments | Low | Data access parts need changes |
| Environment Setup | Manual setup in Colab | Medium | Install required libraries |
| Data Storage | Move to AWS S3 | High | Security, costs, performance |
| Memory Management | Potential adjustments | Medium | Colab may have different limits |
| Runtime | Adjust for longer sessions | Low | Colab offers up to 12-hour sessions |
| Persistent Storage | Implement solution | Medium | Use Google Drive or re-download |
| Version Control | Set up GitHub integration | Medium | Replace Kaggle's built-in system |
| Result Submission | Manual process | Low | Download from Colab, upload to Kaggle |



MODELS REVIEW



BINARY CLASSIFICATION ALGORITHM REVIEW

| | Convolutional Neural Network (CNN) | Transfer Learning with Pre-trained Models (e.g., ResNet, EfficientNet) | Ensemble Methods (e.g., Random Forest, Gradient Boosting) | Support Vector Machines (SVM) with Kernel Trick | Attention-based Models (e.g., Vision Transformer - ViT) |
|-----------------|--|---|--|--|--|
| Characteristics | <ul style="list-style-type: none"> Specialized for image processing Automatically learns relevant features from images Deep architecture allows for complex pattern recognition | <ul style="list-style-type: none"> Utilizes pre-trained models on large datasets Fine-tunes the model for specific task Leverages learned features from diverse image datasets | <ul style="list-style-type: none"> Combines multiple models to improve performance Can use different types of base models Often used with decision trees as base learners | <ul style="list-style-type: none"> Finds optimal hyperplane to separate classes Can use different kernel functions for non-linear separation Focuses on maximizing the margin between classes | <ul style="list-style-type: none"> Uses attention mechanisms to process image patches Can capture long-range dependencies in images Inspired by transformer architectures in NLP |
| Advantages | <ul style="list-style-type: none"> Excellent performance on image classification tasks Can capture spatial hierarchies in images Robust to variations in input | <ul style="list-style-type: none"> Requires less task-specific training data Often achieves high performance quickly Benefits from features learned on large, diverse datasets | <ul style="list-style-type: none"> Generally provides better performance than single models Reduces overfitting Can handle different types of features | <ul style="list-style-type: none"> Effective in high-dimensional spaces Memory efficient Versatile through different kernel functions | <ul style="list-style-type: none"> Can achieve state-of-the-art performance on many vision tasks Scales well with larger datasets and model sizes Can potentially capture more global context than CNNs |
| Disadvantages | <ul style="list-style-type: none"> Requires large amounts of labeled data Computationally intensive Can be prone to overfitting on small datasets | <ul style="list-style-type: none"> May require fine-tuning of hyperparameters Can be computationally expensive for very large models May not be optimal if target task is significantly different from pre-training task | <ul style="list-style-type: none"> May be computationally expensive Can be more complex to interpret Requires careful tuning of ensemble parameters | <ul style="list-style-type: none"> Can be slow to train on large datasets Sensitive to feature scaling May require careful parameter tuning | <ul style="list-style-type: none"> May require large amounts of data for training from scratch Can be computationally expensive May struggle with small objects or fine-grained details |
| Code example | | | | | |

MOST POWERFUL ALGORITHM COMPARISON FOR IMAGE CLASSIFICATION

| | EfficientNet | DenseNet | ResNet | Inception-ResNet | Vision Transformer (ViT) |
|------------------------|---|---|---|---|---|
| Definition | CNN architecture using compound scaling to balance network depth, width, and resolution | CNN with dense connectivity pattern between layers | Deep CNN using residual learning | Combines Inception architecture with residual connections | Applies transformer architecture to image patches |
| Characteristics | <ul style="list-style-type: none"> Compound scaling Mobile Inverted Bottleneck Convolution blocks Squeeze-and-Excitation blocks | <ul style="list-style-type: none"> Dense connectivity Feature reuse Compact architecture | <ul style="list-style-type: none"> Residual learning Skip connections Very deep architecture | <ul style="list-style-type: none"> Multi-scale feature extraction Residual connections Inception modules | <ul style="list-style-type: none"> Patch-based image processing Self-attention mechanism No convolutions |
| Advantages | <ul style="list-style-type: none"> Excellent efficiency-accuracy trade-off Scalable architecture Strong performance on various tasks | <ul style="list-style-type: none"> Efficient parameter usage Strong feature propagation Mitigates vanishing gradient | <ul style="list-style-type: none"> Enables training of very deep networks Widely applicable Well-understood architecture | <ul style="list-style-type: none"> High accuracy Multi-scale feature capture Benefits of both Inception and ResNet | <ul style="list-style-type: none"> Captures global dependencies Scales well to large datasets Potential for cross-modal applications |
| Disadvantages | <ul style="list-style-type: none"> Complex architecture May require large datasets for full benefit | <ul style="list-style-type: none"> Memory intensive during training Complex dense connections | <ul style="list-style-type: none"> Very deep variants can be computationally expensive | <ul style="list-style-type: none"> Computationally expensive Complex architecture | <ul style="list-style-type: none"> Requires large datasets for training from scratch May struggle with small objects |
| Family Models | B0, B1, B2, B3, B4, B5, B6, B7, L2, V2 (S, M, L) | DenseNet-121, 169, 201, 264 | ResNet-18, 34, 50, 101, 152, 50V2, 101V2, 152V2 | Inception-ResNet-v1, Inception-ResNet-v2 | ViT-Base, ViT-Large, ViT-Huge |
| Comparison of Variants | B0 (smallest) to B7 (largest) Each step increases accuracy and complexity V2 improves training speed and accuracy | Deeper variants (201, 264) offer higher accuracy but more parameters 121 often good balance of efficiency and performance | Deeper variants (101, 152) offer higher accuracy V2 variants improve training stability | v2 generally preferred over v1 for better performance | Larger variants (Large, Huge) offer higher accuracy but require more data and compute |

5 SELECTED MODELS FOR SKIN CANCER DETECTION

| | EfficientNet-B7 | ResNet-152V2 | DenseNet-201 | Inception-ResNet-V2 | Vision Transformer (ViT) |
|-----------------|--|---|---|--|--|
| | Compound scaling method to balance network depth, width, and resolution. B7 variant is one of the largest and most powerful | Residual connections, allowing for very deep networks. The 152-layer version with V2 improvements is a powerful model for image classification | Connects each layer to every other layer in a feed-forward fashion, which helps mitigate the vanishing gradient problem and encourage feature reuse | Combines the Inception architecture, which uses multiple filter sizes in each layer, with residual connections from ResNet | Excellent performance on image classification tasks and represents a different paradigm in computer vision (not a traditional CNN) |
| Characteristics | <ul style="list-style-type: none"> Optimized architecture for both accuracy and efficiency Achieved state-of-the-art performance on ImageNet Uses compound scaling to systematically scale network dimensions | <ul style="list-style-type: none"> Very deep architecture with residual connections Improved training stability for deep networks Strong feature extraction capabilities | <ul style="list-style-type: none"> Dense connectivity pattern between layers Requires fewer parameters than traditional CNNs Encourages feature reuse throughout the network | <ul style="list-style-type: none"> Combines Inception modules with residual connections Very deep and wide network Efficient use of computational resources | <ul style="list-style-type: none"> Treats image patches as tokens and applies self-attention Inspired by transformer models in NLP Can capture long-range dependencies in images |
| Advantages | <ul style="list-style-type: none"> Excellent performance-to-parameter ratio Can capture fine-grained details important in medical imaging Scalable architecture allows for different model sizes | <ul style="list-style-type: none"> Proven architecture in many computer vision tasks Can capture complex hierarchical features Good trade-off between depth and computational efficiency | <ul style="list-style-type: none"> Efficient parameter usage Strong gradient flow throughout the network Can capture fine-grained features effectively | <ul style="list-style-type: none"> High accuracy on image classification tasks Can capture features at multiple scales simultaneously Benefits from both Inception and ResNet architectures | <ul style="list-style-type: none"> State-of-the-art performance on many vision tasks Can potentially capture more global context than CNNs Scales well with larger datasets and model sizes |

CODE EXAMPLE

```
# EfficientNet-B7 model
from tensorflow.keras.applications import EfficientNetB7

from tensorflow.keras.layers import GlobalAveragePooling2D, Dense
from tensorflow.keras.models import Model

base_model = EfficientNetB7(weights='imagenet', include_top=False,
input_shape=(224, 224, 3))

x = GlobalAveragePooling2D()(base_model.output)
x = Dense(256, activation='relu')(x)
output = Dense(1, activation='sigmoid')(x)

model = Model(inputs=base_model.input, outputs=output)

for layer in base_model.layers:
    layer.trainable = False

model.compile(optimizer='adam', loss='binary_crossentropy',
metrics=['accuracy'])
```

```
# Vision Transformer (ViT) model
import timm

import torch

import torch.nn as nn

class ViTForBinaryClassification(nn.Module):

    def __init__(self, num_classes=1):

        super().__init__()

        self.vit = timm.create_model('vit_base_patch16_224',
pretrained=True)

        self.vit.head = nn.Linear(self.vit.head.in_features,
num_classes)

    def forward(self, x):

        return torch.sigmoid(self.vit(x))

model = ViTForBinaryClassification()
criterion = nn.BCELoss()

optimizer = torch.optim.Adam(model.parameters(), lr=1e-4)
```



SCIENTIFIC REPORT ABOUT MELANOMA DIAGNOSIS



SUMMARY OF THE OPTIMIZED DEEP-CNN ARCHITECTURE WITH CUSTOM MINI-BATCH LOGIC AND LOSS FUNCTION

Source: <https://www.nature.com/articles/s41598-021-96707-8>

- Topic: binary melanoma classification system using deep learning that outperforms 157 dermatologists
- Methodology:
 - Model architecture:
 - An optimized CNN architecture based on DenseNet169
 - Retrained of fully connected layers
 - Model innovation:
 - Custom Loss Function to handle imbalanced data
 - Custom mini-batch logic to maintain a fixed ratio between classes
 - Real-time data augmentation
 - Optimization: Adam optimizer & implementation of cyclical learning rate
 - Experimentation: Comparison of three models: ORI (original), BON (with custom batch), BLF (with custom batch and loss function) on ISIC 2019 test set (Test-10) and MClass-D dataset
- Main results:
 - The BLF model achieves an AUC of 94.4%
 - Sensitivity of 85.0% and specificity of 95.0% with a prediction threshold of 0.5
 - Balanced performance: 90.0% sensitivity and 93.8% specificity with an adjusted threshold
 - Outperforms all 157 tested dermatologists on MClass-D

SUMMARY OF THE OPTIMIZED DEEP-CNN ARCHITECTURE WITH CUSTOM MINI-BATCH LOGIC AND LOSS FUNCTION

Source: <https://www.nature.com/articles/s41598-021-96707-8>

- Performance analysis:
 - The custom loss function improves the balance between sensitivity and specificity
 - Custom mini-batch logic improves training stability
 - DenseNet169 architecture proves more effective than InceptionV3 and ResNet50 for this task
- Implications and perspectives:
 - Potential for application in computer-assisted medical diagnosis
 - Possibility to extend the approach to other medical image classification tasks
 - Need for further research on custom loss functions and fully connected layer architecture
- Limitations:
 - Need for validation on external datasets and in real clinical conditions
 - Necessity to interpret model results for medical use
- Contribution to the state of the art:
 - First approach outperforming all tested dermatologists on the MClass-D dataset
 - New method for handling imbalanced datasets in medical image classification